

Warrant-bearing instance

What licenses the target language?

Model-visible workflow input

Selected warrant field
licenses the target language

Evidence span

Distractor language cues

Output contract

Authorized target language

Integrity gates

Which threats are
blocked before scoring?

Schema validity

Span-grounded warrant

Triad invariants

Shortcut-resistance checks

Contract alignment

blocks broken warrants, invalid triads,
shortcut-only labels, and version drift

Validity-preserved instance

Fixed scoring instrument

How are raw outputs
converted into metrics?

Immutable raw model output

Visible response region

Predicted response language

Validity guards: task, format,
field policy

LWA, VRS, TWCS, TVRS

Measured model behavior

Audit-bounded reporting

How are repairs and
audits interpreted?

Narrow scorer repair

Regression preservation

Human audit addenda

Claim boundary

**Audits bound claims,
not task definition.**

Bounded benchmark claim